

# prmon Memory Anomaly Detection

## Warm-up Project

Agnidipto Banik

February 28, 2026

# Objective

- Generate memory time-series data using prmon
- Inject a controlled memory anomaly using burner tests
- Apply statistical anomaly detection
- Evaluate detection performance quantitatively and visually

# Experimental Setup

- Environment: WSL Ubuntu + prmon (built from source)
- Metric used: RSS (Resident Set Size)
- Three segments:
  - Normal baseline
  - Elevated memory usage (anomaly)
  - Recovery to baseline
- Total samples: 228

# Anomaly Injection

- Normal burner run establishes baseline behavior
- Memory parameters increased to create sustained elevation
- System returned to baseline to simulate recovery
- Anomaly region labeled using known injection interval

# Data Generation Using prmon

```
# Baseline run
```

```
~/prmon/build/package/prmon \
--interval 1 \
--filename data/normal_part.csv \
-- ~/prmon/build/package/tests/mem-burner \
--malloc 500 --sleep 150
```

```
# Elevated memory (anomaly)
```

```
~/prmon/build/package/prmon \
--interval 1 \
--filename data/anomaly_part.csv \
-- ~/prmon/build/package/tests/mem-burner \
--malloc 650 --sleep 150
```

```
# Recovery run
```

```
~/prmon/build/package/prmon \
--interval 1 \
--filename data/segment3_recovery.csv \
-- ~/prmon/build/package/tests/mem-burner
```

# Dataset Construction

```
s1 = pd.read_csv("../data/normal_part.csv", sep=r"\s+")
s2 = pd.read_csv("../data/anomaly_part.csv", sep=r"\s+
")
s3 = pd.read_csv("../data/segment3_recovery.csv", sep=
r"\s+")

df = pd.concat([s1, s2, s3], ignore_index=True)

df["is_anomaly"] = 0
df.loc[len(s1):len(s1)+len(s2)-1, "is_anomaly"] = 1
```

- Concatenates normal, anomaly, and recovery segments
- Ground truth labels created using known injection interval
- Enables quantitative evaluation

# Rolling Z-Score Implementation

```
window = 30
threshold = 3

df["rolling_mean"] = df["rss"].rolling(window=window).
    mean()
df["rolling_std"] = df["rss"].rolling(window=window).
    std()

df["z_rolling"] = (
    (df["rss"] - df["rolling_mean"]) /
    df["rolling_std"].replace(0, np.nan)
).fillna(0)

df["rolling_detected"] = (df["z_rolling"].abs() >
    threshold).astype(int)
```

- Computes rolling mean and standard deviation
- Standardizes deviation using Z-score
- Flags  $|z| > 3$  as anomaly

# Frozen Baseline Implementation

```
baseline_mean = s1["rss"].mean()  
baseline_std = s1["rss"].std()  
  
df["z_frozen"] = (df["rss"] - baseline_mean) /  
    baseline_std  
df["frozen_detected"] = (df["z_frozen"].abs() >  
    threshold).astype(int)
```

- Baseline statistics computed from clean segment
- Fixed reference used for all future samples
- Detects sustained structural shifts
- Does not adapt to new regimes

# Evaluation Metrics

```
precision_score(df["is_anomaly"], df[col])
recall_score(df["is_anomaly"], df[col])
f1_score(df["is_anomaly"], df[col])
confusion_matrix(df["is_anomaly"], df[col])
```

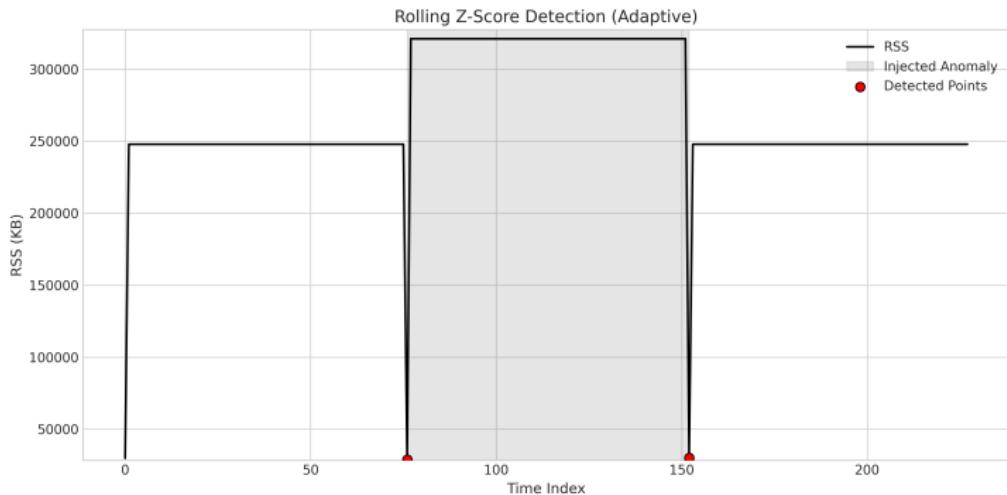
- Precision: false positive control
- Recall: anomaly coverage
- F1-score: balance between precision and recall
- Confusion matrix provides detailed breakdown

# Rolling Z-Score (Adaptive Detection)

$$z_t = \frac{x_t - \mu_{window}}{\sigma_{window}}$$

- Compares current value to recent behavior
- Adapts dynamically to new levels
- Threshold:  $|z| > 3$

# Rolling Z-Score – Results



# Why Rolling Fails

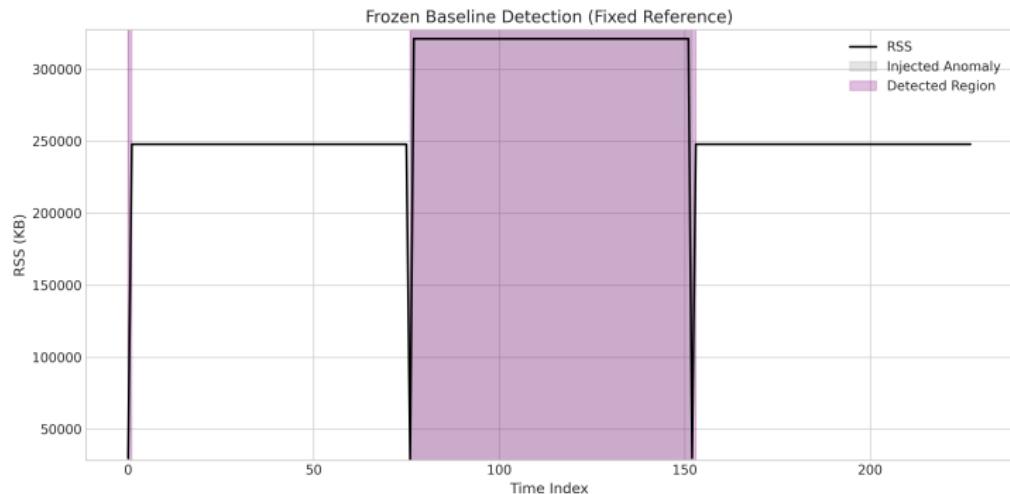
- Detects boundary transitions only
- Rolling window adapts to sustained shift
- Elevated regime becomes statistically normal
- Low recall for step anomalies

# Frozen Baseline Z-Score (Fixed Reference)

$$z_t = \frac{x_t - \mu_{baseline}}{\sigma_{baseline}}$$

- Baseline statistics computed from initial normal segment
- No adaptation to new regimes
- Detects sustained structural deviation

# Frozen Baseline – Results



# Why Frozen Performs Better

- Step anomaly introduces sustained mean shift
- Frozen model remains anchored to healthy baseline
- Entire elevated regime remains statistically significant
- High precision and recall observed

# Final Quantitative Results

## Rolling Z-Score

$$\begin{bmatrix} 151 & 1 \\ 75 & 1 \end{bmatrix}$$

## Frozen Baseline

$$\begin{bmatrix} 150 & 2 \\ 0 & 76 \end{bmatrix}$$

# Why Not Machine Learning?

- Warm-up focused on controlled statistical evaluation
- Dataset is small and structured
- Statistical methods are interpretable and sufficient

# Trade-offs and Sustainability

- Rolling suitable for short-lived spikes
- Frozen requires clean baseline period
- Frozen may struggle with gradual drift
- Real deployments require adaptive baselines and multi-metric monitoring

# Conclusion

- Controlled memory anomaly successfully injected
- Statistical detection methods applied and evaluated
- Frozen baseline outperformed rolling detection
- Demonstrates structured experimental reasoning